

# fraud-detection

February 21, 2024

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
```

```
[2]: data = pd.read_csv('/content/creditcard.csv')
data.head()
```

```
[2]:
```

	Time	V1	V2	V3	V4	V5	V6	V7 \
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941

	V8	V9	...	V21	V22	V23	V24	V25 \
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376
4	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010

	V26	V27	V28	Amount	Class
0	-0.189115	0.133558	-0.021053	149.62	0.0
1	0.125895	-0.008983	0.014724	2.69	0.0
2	-0.139097	-0.055353	-0.059752	378.66	0.0
3	-0.221929	0.062723	0.061458	123.50	0.0
4	0.502292	0.219422	0.215153	69.99	0.0

[5 rows x 31 columns]

```
[3]: data.describe()
```

```
[3]:
```

	Time	V1	V2	V3 \
count	182329.000000	182329.000000	182329.000000	182329.000000
mean	64969.194676	-0.140528	0.022185	0.417493
std	30828.520647	1.873983	1.614571	1.416903

min	0.000000	-56.407510	-72.715728	-33.680984
25%	42410.000000	-0.975751	-0.560485	-0.172141
50%	63339.000000	-0.139500	0.094318	0.570588
75%	81849.000000	1.196395	0.791998	1.267890
max	125353.000000	2.439207	22.057729	9.382558

	V4	V5	V6	V7 \
count	182329.000000	182329.000000	182329.000000	182329.000000
mean	0.096426	-0.150145	0.053535	-0.068736
std	1.384183	1.355202	1.304002	1.212992
min	-5.683171	-42.147898	-26.160506	-43.557242
25%	-0.768971	-0.809987	-0.699692	-0.582837
50%	0.099818	-0.203398	-0.209359	-0.023584
75%	0.914961	0.412695	0.447010	0.477721
max	16.875344	34.801666	22.529298	36.677268

	V8	V9	...	V21	V22 \
count	182329.000000	182328.000000	...	182328.000000	182328.000000
mean	0.025907	0.014692	...	-0.021476	-0.067798
std	1.223914	1.144658	...	0.742347	0.678239
min	-73.216718	-13.434066	...	-34.830382	-10.933144
25%	-0.170715	-0.661010	...	-0.228923	-0.542896
50%	0.050375	-0.076790	...	-0.049453	-0.053276
75%	0.344873	0.634247	...	0.138317	0.389549
max	20.007208	15.594995	...	27.202839	10.503090

	V23	V24	V25	V26 \
count	182328.000000	182328.000000	182328.000000	182328.000000
mean	-0.019683	0.007206	0.078579	0.010185
std	0.594938	0.600877	0.476265	0.489616
min	-44.807735	-2.836627	-10.295397	-2.604551
25%	-0.169356	-0.335944	-0.216393	-0.330644
50%	-0.032859	0.056785	0.122284	-0.059203
75%	0.106009	0.418491	0.393032	0.268662
max	19.002942	4.022866	7.519589	3.517346

	V27	V28	Amount	Class
count	182328.000000	182328.000000	182328.000000	182328.000000
mean	0.002032	0.002194	88.52699	0.002002
std	0.392343	0.307251	247.55138	0.044698
min	-22.565679	-11.710896	0.00000	0.000000
25%	-0.066026	-0.034975	5.76000	0.000000
50%	0.007456	0.020019	22.43000	0.000000
75%	0.089507	0.078215	78.12000	0.000000
max	12.152401	33.847808	19656.53000	1.000000

[8 rows x 31 columns]

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 182329 entries, 0 to 182328
Data columns (total 31 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   Time    182329 non-null  float64
 1   V1      182329 non-null  float64
 2   V2      182329 non-null  float64
 3   V3      182329 non-null  float64
 4   V4      182329 non-null  float64
 5   V5      182329 non-null  float64
 6   V6      182329 non-null  float64
 7   V7      182329 non-null  float64
 8   V8      182329 non-null  float64
 9   V9      182328 non-null  float64
10  V10     182328 non-null  float64
11  V11     182328 non-null  float64
12  V12     182328 non-null  float64
13  V13     182328 non-null  float64
14  V14     182328 non-null  float64
15  V15     182328 non-null  float64
16  V16     182328 non-null  float64
17  V17     182328 non-null  float64
18  V18     182328 non-null  float64
19  V19     182328 non-null  float64
20  V20     182328 non-null  float64
21  V21     182328 non-null  float64
22  V22     182328 non-null  float64
23  V23     182328 non-null  float64
24  V24     182328 non-null  float64
25  V25     182328 non-null  float64
26  V26     182328 non-null  float64
27  V27     182328 non-null  float64
28  V28     182328 non-null  float64
29  Amount  182328 non-null  float64
30  Class   182328 non-null  float64
dtypes: float64(31)
memory usage: 43.1 MB
```

```
[5]: data.isna().sum()
```

```
[5]: Time      0
     V1        0
     V2        0
     V3        0
```

V4	0
V5	0
V6	0
V7	0
V8	0
V9	1
V10	1
V11	1
V12	1
V13	1
V14	1
V15	1
V16	1
V17	1
V18	1
V19	1
V20	1
V21	1
V22	1
V23	1
V24	1
V25	1
V26	1
V27	1
V28	1
Amount	1
Class	1

dtype: int64

### Data Preprocessing:

```
[6]: from sklearn.impute import SimpleImputer

# Impute missing values with the mean for numerical columns
imputer = SimpleImputer(strategy="mean")

# Fill missing values in numerical columns
data = pd.DataFrame(imputer.fit_transform(data), columns=data.columns)

# Check if any missing values remain
print(data.isnull().sum())
```

Time	0
V1	0
V2	0
V3	0
V4	0
V5	0

```
V6          0
V7          0
V8          0
V9          0
V10         0
V11         0
V12         0
V13         0
V14         0
V15         0
V16         0
V17         0
V18         0
V19         0
V20         0
V21         0
V22         0
V23         0
V24         0
V25         0
V26         0
V27         0
V28         0
Amount      0
Class       0
dtype: int64
```

```
[7]: from sklearn.model_selection import train_test_split

# Separate features (X) and target variable (y)
X = data.drop(columns=["Class"]) # Features
y = data["Class"] # Target variable

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↪random_state=42)
```

### Feature Selection:

```
[9]: from sklearn.preprocessing import StandardScaler

# Create a StandardScaler object
scaler = StandardScaler()

# Fit the scaler to the training data and transform it
X_train_scaled = scaler.fit_transform(X_train)

# Transform the testing data using the fitted scaler
```

```
X_test_scaled = scaler.transform(X_test)
```

### Data Splitting:

```
[10]: from sklearn.model_selection import train_test_split

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↪random_state=42)
```

### Model Selection:

```
[11]: from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score,
↪f1_score, roc_auc_score
```

```
[12]: # Define models
models = {
    'Logistic Regression': LogisticRegression(max_iter=1000),
    'Random Forest': RandomForestClassifier(),
    'Gradient Boosting': GradientBoostingClassifier(),
    'Neural Network': MLPClassifier(max_iter=1000) # adjusting max_iter based
↪on convergence
}
```

```
[13]: # Define a threshold value to classify as 1 (fraudulent)
threshold_value = 0.5

# Convert continuous target values to binary labels
y_train = (y_train > threshold_value).astype(int)
```

### Model Evaluation:

```
[14]: # Train and evaluate models
for name, model in models.items():
    print(f"\033[1mTraining {name}:\033[0m")
    model.fit(X_train, y_train)

    # Predict on the test set
    y_pred = model.predict(X_test)

    # Evaluate the model
    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
```

```

auc_roc = roc_auc_score(y_test, y_pred)

# Print evaluation metrics
print(f"\033[1;3mAccuracy:\033[0m {accuracy:.4f}, \n\033[1;3mPrecision:
↪\033[0m {precision:.4f}, \n\033[1;3mRecall:\033[0m {recall:.4f}, \n\033[1;
↪3mF1-score:\033[0m {f1:.4f}, \n\033[1;3mAUC-ROC:\033[0m {auc_roc:.4f}\n")

```

#### Training Logistic Regression:

Accuracy: 0.9990,  
 Precision: 0.8085,  
 Recall: 0.5758,  
 F1-score: 0.6726,  
 AUC-ROC: 0.7878

#### Training Random Forest:

Accuracy: 0.9995,  
 Precision: 0.9259,  
 Recall: 0.7576,  
 F1-score: 0.8333,  
 AUC-ROC: 0.8787

#### Training Gradient Boosting:

Accuracy: 0.9989,  
 Precision: 0.8000,  
 Recall: 0.5455,  
 F1-score: 0.6486,  
 AUC-ROC: 0.7726

#### Training Neural Network:

Accuracy: 0.9979,  
 Precision: 0.4479,  
 Recall: 0.6515,  
 F1-score: 0.5309,  
 AUC-ROC: 0.8250

#### Data Visualization:

```

[15]: from sklearn.metrics import confusion_matrix, roc_curve,
↪precision_recall_curve, auc


```

```

[16]: # Plot confusion matrix
y_pred = model.predict(X_test)
conf_matrix = confusion_matrix(y_test, y_pred)
plt.imshow(conf_matrix, cmap='Blues', interpolation='nearest')
plt.title('Confusion Matrix')
plt.colorbar()
plt.xlabel('Predicted Label')


```

```
plt.ylabel('True Label')
plt.show()
```



```
[17]: # Plot ROC curve
y_pred_proba = model.predict_proba(X_test)[: , 1]
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (AUC = {roc_auc:.
    2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```





```
[18]: # Plot precision-recall curve
precision, recall, _ = precision_recall_curve(y_test, y_pred_proba)
plt.plot(recall, precision, color='blue', lw=2, label='Precision-Recall curve')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.legend(loc='upper right')
plt.show()
```

