

housing-price-prediction

February 19, 2024

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
```

```
[2]: df = pd.read_csv("/content/Housing.csv")
df.head()
```

```
[2]:      price  area  bedrooms  bathrooms  stories  mainroad  guestroom  basement  \
0  13300000  7420         4           2         3        yes         no         no
1  12250000  8960         4           4         4        yes         no         no
2  12250000  9960         3           2         2        yes         no         yes
3  12215000  7500         4           2         2        yes         no         yes
4  11410000  7420         4           1         2        yes         yes        yes

      hotwaterheating  airconditioning  parking  prefarea  furnishingstatus
0                no                yes         2        yes        furnished
1                no                yes         3         no        furnished
2                no                no         2        yes    semi-furnished
3                no                yes         3        yes        furnished
4                no                yes         2         no        furnished
```

Data Cleaning:

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                 545 non-null   int64
1   area                 545 non-null   int64
2   bedrooms             545 non-null   int64
3   bathrooms            545 non-null   int64
4   stories              545 non-null   int64
5   mainroad             545 non-null   object
6   guestroom            545 non-null   object
7   basement             545 non-null   object
```

```

8   hotwaterheating    545 non-null    object
9   airconditioning    545 non-null    object
10  parking            545 non-null    int64
11  prefarea           545 non-null    object
12  furnishingstatus   545 non-null    object
dtypes: int64(6), object(7)
memory usage: 55.5+ KB

```

```
[4]: df.describe()
```

```

[4]:
      price      area  bedrooms  bathrooms  stories \
count  5.450000e+02   545.000000  545.000000  545.000000  545.000000
mean   4.766729e+06   5150.541284    2.965138    1.286239    1.805505
std    1.870440e+06   2170.141023    0.738064    0.502470    0.867492
min    1.750000e+06   1650.000000    1.000000    1.000000    1.000000
25%    3.430000e+06   3600.000000    2.000000    1.000000    1.000000
50%    4.340000e+06   4600.000000    3.000000    1.000000    2.000000
75%    5.740000e+06   6360.000000    3.000000    2.000000    2.000000
max    1.330000e+07  16200.000000    6.000000    4.000000    4.000000

      parking
count  545.000000
mean    0.693578
std     0.861586
min     0.000000
25%     0.000000
50%     0.000000
75%     1.000000
max     3.000000

```

```

[5]: print("\033[1mMissing values:\033[0m")
      print(df.isnull().sum())

```

```

Missing values:
price      0
area       0
bedrooms   0
bathrooms  0
stories    0
mainroad   0
guestroom  0
basement   0
hotwaterheating  0
airconditioning  0
parking     0
prefarea    0
furnishingstatus  0
dtype: int64

```

```
[6]: # Convert categorical variables to numerical using one-hot encoding
df = pd.get_dummies(df, columns=['mainroad', 'guestroom', 'basement',
    ↪ 'hotwaterheating',
    'airconditioning', 'prefarea',
    ↪ 'furnishingstatus'], drop_first=True)
```

```
[7]: # Verify changes
print("\033[1mDataFrame after cleaning:\033[0m")
df.head()
```

DataFrame after cleaning:

```
[7]:
```

	price	area	bedrooms	bathrooms	stories	parking	mainroad_yes \
0	13300000	7420	4	2	3	2	1
1	12250000	8960	4	4	4	3	1
2	12250000	9960	3	2	2	2	1
3	12215000	7500	4	2	2	3	1
4	11410000	7420	4	1	2	2	1

	guestroom_yes	basement_yes	hotwaterheating_yes	airconditioning_yes \
0	0	0	0	1
1	0	0	0	1
2	0	1	0	0
3	0	1	0	1
4	1	1	0	1

	prefarea_yes	furnishingstatus_semi-furnished	furnishingstatus_unfurnished
0	1	0	0
1	0	0	0
2	1	1	0
3	1	0	0
4	0	0	0

Feature Selection:

```
[8]: df.corr()
```

```
[8]:
```

	price	area	bedrooms	bathrooms \
price	1.000000	0.535997	0.366494	0.517545
area	0.535997	1.000000	0.151858	0.193820
bedrooms	0.366494	0.151858	1.000000	0.373930
bathrooms	0.517545	0.193820	0.373930	1.000000
stories	0.420712	0.083996	0.408564	0.326165
parking	0.384394	0.352980	0.139270	0.177496
mainroad_yes	0.296898	0.288874	-0.012033	0.042398
guestroom_yes	0.255517	0.140297	0.080549	0.126469
basement_yes	0.187057	0.047417	0.097312	0.102106
hotwaterheating_yes	0.093073	-0.009229	0.046049	0.067159

airconditioning_yes	0.452954	0.222393	0.160603	0.186915
prefarea_yes	0.329777	0.234779	0.079023	0.063472
furnishingstatus_semi-furnished	0.063656	0.006156	0.050040	0.029834
furnishingstatus_unfurnished	-0.280587	-0.142278	-0.126252	-0.132107

	stories	parking	mainroad_yes	\
price	0.420712	0.384394	0.296898	
area	0.083996	0.352980	0.288874	
bedrooms	0.408564	0.139270	-0.012033	
bathrooms	0.326165	0.177496	0.042398	
stories	1.000000	0.045547	0.121706	
parking	0.045547	1.000000	0.204433	
mainroad_yes	0.121706	0.204433	1.000000	
guestroom_yes	0.043538	0.037466	0.092337	
basement_yes	-0.172394	0.051497	0.044002	
hotwaterheating_yes	0.018847	0.067864	-0.011781	
airconditioning_yes	0.293602	0.159173	0.105423	
prefarea_yes	0.044425	0.091627	0.199876	
furnishingstatus_semi-furnished	-0.003648	0.041327	0.011450	
furnishingstatus_unfurnished	-0.082972	-0.165705	-0.133123	

	guestroom_yes	basement_yes	\
price	0.255517	0.187057	
area	0.140297	0.047417	
bedrooms	0.080549	0.097312	
bathrooms	0.126469	0.102106	
stories	0.043538	-0.172394	
parking	0.037466	0.051497	
mainroad_yes	0.092337	0.044002	
guestroom_yes	1.000000	0.372066	
basement_yes	0.372066	1.000000	
hotwaterheating_yes	-0.010308	0.004385	
airconditioning_yes	0.138179	0.047341	
prefarea_yes	0.160897	0.228083	
furnishingstatus_semi-furnished	0.005821	0.050284	
furnishingstatus_unfurnished	-0.099023	-0.117935	

	hotwaterheating_yes	airconditioning_yes	\
price	0.093073	0.452954	
area	-0.009229	0.222393	
bedrooms	0.046049	0.160603	
bathrooms	0.067159	0.186915	
stories	0.018847	0.293602	
parking	0.067864	0.159173	
mainroad_yes	-0.011781	0.105423	
guestroom_yes	-0.010308	0.138179	
basement_yes	0.004385	0.047341	

hotwaterheating_yes	1.000000	-0.130023
airconditioning_yes	-0.130023	1.000000
prefarea_yes	-0.059411	0.117382
furnishingstatus_semi-furnished	0.063819	-0.053179
furnishingstatus_unfurnished	-0.059194	-0.094086

	prefarea_yes \
price	0.329777
area	0.234779
bedrooms	0.079023
bathrooms	0.063472
stories	0.044425
parking	0.091627
mainroad_yes	0.199876
guestroom_yes	0.160897
basement_yes	0.228083
hotwaterheating_yes	-0.059411
airconditioning_yes	0.117382
prefarea_yes	1.000000
furnishingstatus_semi-furnished	-0.011535
furnishingstatus_unfurnished	-0.081271

	furnishingstatus_semi-furnished \
price	0.063656
area	0.006156
bedrooms	0.050040
bathrooms	0.029834
stories	-0.003648
parking	0.041327
mainroad_yes	0.011450
guestroom_yes	0.005821
basement_yes	0.050284
hotwaterheating_yes	0.063819
airconditioning_yes	-0.053179
prefarea_yes	-0.011535
furnishingstatus_semi-furnished	1.000000
furnishingstatus_unfurnished	-0.588405

	furnishingstatus_unfurnished
price	-0.280587
area	-0.142278
bedrooms	-0.126252
bathrooms	-0.132107
stories	-0.082972
parking	-0.165705
mainroad_yes	-0.133123
guestroom_yes	-0.099023


basement_yes	-0.117935
hotwaterheating_yes	-0.059194
airconditioning_yes	-0.094086
prefarea_yes	-0.081271
furnishingstatus_semi-furnished	-0.588405
furnishingstatus_unfurnished	1.000000

```
[9]: import matplotlib.pyplot as plt
import seaborn as sns

corr = df.corr()

fig, ax = plt.subplots(figsize=(8, 8))
sns.heatmap(corr, annot=True, linewidths=0.2, linecolor="black", fmt=".2f",
            ↪ax=ax)

plt.show()
```



```
[10]: # Find features with correlation greater than a threshold with the target
      ↪variable
      threshold = 0.5
      high_corr_features = corr.index[abs(corr['price']) > threshold].tolist()

      # Print selected features
      print("\033[1mFeatures with high correlation with the target variable:\033[0m")
      print(high_corr_features)
```

Features with high correlation with the target variable:
['price', 'area', 'bathrooms']

```
[11]: from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression

# Assuming X contains the features and y contains the target variable (price)
X = df.drop(columns=['price']) # Features
y = df['price'] # Target variable

# Initialize a linear regression model
model = LinearRegression()

# Initialize RFE
rfe = RFE(model, n_features_to_select=3)
# Fit RFE
rfe.fit(X, y)

# Get selected features
selected_features = X.columns[rfe.support_]
# Print selected features
print("\033[1mSelected features using RFE:\033[0m")
print(selected_features)
```

Selected features using RFE:

Index(['bathrooms', 'mainroad_yes', 'airconditioning_yes'], dtype='object')

```
[12]: from sklearn.linear_model import Lasso
from sklearn.preprocessing import StandardScaler

# Assuming X contains the features and y contains the target variable (price)
X = df.drop(columns=['price']) # Features
y = df['price'] # Target variable

# Standardize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Initialize Lasso regression model
lasso = Lasso(alpha=0.1)

# Fit Lasso model
lasso.fit(X_scaled, y)

# Get coefficients and select non-zero coefficient features
selected_features = X.columns[lasso.coef_ != 0]

# Print selected features
print("\033[1mSelected features using Lasso regularization:\033[0m")
print(selected_features)
```


Selected features using Lasso regularization:

```
Index(['area', 'bedrooms', 'bathrooms', 'stories', 'parking', 'mainroad_yes',  
      'guestroom_yes', 'basement_yes', 'hotwaterheating_yes',  
      'airconditioning_yes', 'prefarea_yes',  
      'furnishingstatus_semi-furnished', 'furnishingstatus_unfurnished'],  
      dtype='object')
```

Model Training:

```
[13]: from sklearn.linear_model import LinearRegression  
  
# Assuming X contains the selected features and y contains the target variable_  
↳(price)  
X = df[selected_features] # Selected features  
y = df['price'] # Target variable  
  
# Initialize and fit the linear regression model  
model = LinearRegression()  
model.fit(X, y)
```

```
[13]: LinearRegression()
```

Model Evaluation:

```
[14]: from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error  
from sklearn.model_selection import train_test_split  
  
# Assuming X_test contains the selected features and y_test contains the actual_  
↳target variable values for the test set  
X_test = df[selected_features] # Selected features from the test set  
y_test = df['price'] # Actual target variable values for the test set  
  
# Make predictions on the test set  
y_pred = model.predict(X_test)  
  
# Calculate Mean Squared Error (MSE)  
mse = mean_squared_error(y_test, y_pred)  
  
# Calculate R-squared  
r_squared = r2_score(y_test, y_pred)  
  
# Calculate Mean Absolute Error (MAE)  
mae = mean_absolute_error(y_test, y_pred)  
  
print("\033[1mMean Squared Error (MSE):\033[0m", mse)  
print("\033[1mR-squared:\033[0m", r_squared)  
print("\033[1mMean Absolute Error (MAE):\033[0m", mae)
```

Mean Squared Error (MSE): 1111187722284.4001


R-squared: 0.6818018485540142

Mean Absolute Error (MAE): 775054.3287400283


Visualization:

```
[15]: # price vs area
plt.figure(figsize=(6, 4))
df.plot(kind='scatter', x='price', y='area', s=32, alpha=.8)
plt.gca().spines[['top', 'right']].set_visible(False)
```

<Figure size 600x400 with 0 Axes>



```
[16]: # Plotting the predicted vs. actual values
plt.figure(figsize=(6, 4))
plt.scatter(y_test, y_pred, color='blue', label='Actual vs. Predicted')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         linestyle='--', label='Perfect Prediction')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Actual vs. Predicted Values')
plt.legend()
plt.show()
```



```
[17]: # Plotting residuals
residuals = y_test - y_pred
plt.figure(figsize=(6, 4))
plt.scatter(y_pred, residuals, color='green')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.title('Residuals Plot')
plt.axhline(y=0, color='red', linestyle='--')
plt.show()
```

